

Layouts, Widgets, and Stacking Order

Introduction

- In the previous lecture we discussed the event driven model
 - As the basis of all GUI program development.
 - In this lecture we will now discuss mechanisms for organising our GUI display, presenting information to the user and receiving events from our user
 - Before we cover that however, we will continue with some of the theory behind GUI toolkits namely the stacking order of windows

Introduction

- At the moment we have a very good understanding of the event loop and its role in the maintenance of GUI applications
 - However, we need more than the event loop we need to have mechanisms for displaying controls to the user and receiving their input.
 - The event model is useless, unless there is something to generate events and inputs to that model.
 - We also need mechanisms for organising these controls
 - Not every display will be of the same shape and size.
 - Users will determine how much space your application will take on their screen at all times
 - Thus if the window size changes, your application must change with it.

Introduction

- In this lecture we will explore how these problems can be solved in GUI toolkits.
 - We will come across the stacking order for determining which windows are ontop and on the bottom
 - Will also be used to determine the owners of events.
 - We will also discuss widgets/controls
 - To display information to users and also to receive their input.
 - Here is where our event handling will be exposed

Introduction

- In this lecture we will explore how these problems can be solved in GUI toolkits.
 - Finally we will discuss layouts
 - Will be used to determine how widgets get space should the window be resized.

Stacking Order

- Lets assume that we have created a window manager that can generate GUI components.
 - However, it can only show windows side by side but never on top.
 - What if there was a way that we can organise windows in such that we can make them overlap
 - Well one such method is to use a combination of a window stack and painters algorithm to render multiple overlapping windows
 - However, the stack that is used here is not a normal stack datastructure

Stacking Order

- A normal stack is a LIFO data structure
 - Elements can only be added or removed from the top
 - In the stack for a window manager elements must be removable from anywhere in the stack so they can be placed on the top.
 - The window that has focus and is active (i.e. the user's attention) should be at the top of the stack at all times.
 - All windows maintain their positions and bounds on the display.
 - All windows that were above this window will subsequently move down one place in the stack.

Stacking Order

- That takes care of the modelling of overlapping windows
 - The rendering is a different issue.
 - How do you render a stack of overlapping windows?

Stacking Order

- You render from bottom of stack to top of stack
 - One window at a time
 - This is known as painter's algorithm
 - The analogy here is that a painter will draw the background first, then the middle ground and finally the foreground
 - All on a 2D canvas
 - Thus these two modifications give us a stacking window manager that is capable of rendering and displaying overlapping windows

Event Handling

- At the moment though
 - There is no process for determining what window owns any given event
 - Fortunately our stacking order will prove useful here
 - Along with bounding box tests.
 - However the logic for handling keyboard events will be different from mouse events
 - Implementing this logic will make a more useful window manager

Keyboard Event Handling

- Keyboard event handling is simple
 - Look for the window that has current focus
 - Usually the window on top of the stack
 - But depends on whether the window will accept focus
 - And the window manager's policy for focus with new windows.
 - Whichever window has focus will receive all keyboard input
 - If the window does not handle the key event then the window manager may discard it or respond to it
 - Again these are rules decided by the window manager

Mouse Event Handling

- Mouse event handling on the other hand is a bit more difficult
 - We now have a pointer that can be anywhere on the entire screen.
 - Thus it is entirely possible that the user may click on a window that is not on top of the stack
 - Or may not click on a window at all.
 - How do we go about handling these cases?

Mouse Event Handling

- In the first case we have
 - Determine the global coordinates on screen where the user clicked.
 - We then perform a bounding box test on each window in the stack to see which window encloses the click
 - Test from top of stack down to bottom.
 - The first window that contains the click is the window that will own that click event.
 - Thus the click event will be sent to the event handlers of that window and the window will move to top of the stack

Mouse Event Handling

- That only handles the first case
 - However what happens if no window owns that click
 - i.e. the click does not fall in the bounds of any window in the stacking order.
 - It is up to the window manager to then decide what to do with that click
 - Most window managers will interpret a click outside of windows in some way
 - e.g. right clicks will bring up context menus

Event Handling

- So at this point we now have a window manager that can handle overlapping windows
 - And can also handle keyboard and mouse events.
 - That takes care of top level windows however we need to have mechanisms in place that will determine what widget in a GUI window will own a particular event
 - How do we go about doing this?

GUI Widget event handling

- In a GUI application we have a tree data structure called the visual tree.
 - It is made up of all layouts and widgets that belong to a top level window.
 - The application will perform tests to see which widget owns the event
 - Will work from the root of the tree to the leaves.
 - When the widget is found that widget will be given the event by calling it's appropriate event handler
 - The widget will process this event, update itself and if necessary generate and fire more events in response.

GUI Widget event handling

- For example if a button is clicked
 - The mouse click event handler is called
 - The button registers this click
 - Redraws itself on screen
 - Then fires an `ActionEvent` to an attached `EventHandler` so the GUI can respond to that click
 - It will then notify the GUI that the event has been consumed and requires no further processing.

GUI Widget event handling

- This is the full complete cycle of how a window manager gets an event, sends to a GUI application and enables a response to that event
 - Widgets take basic interaction through keyboard and mouse and provide context to those interactions.
- We now have a fully working window manager with event targeting for individual widgets
 - All that is left to sort is how to size and place widgets in the window

Layouts

- How would we do this?

Layouts

- A naive first solution would be to give each widget a position and a size in the window
 - Only works with fixed sized windows
 - Increase in display pixel density will cause the app to suffer.
 - Incapable of scaling up or down
 - No method to allocate extra space.
 - Thus we need some way to determine where widgets are placed and how they are sized
 - How do we do this?

Layouts

- Through a combination of rules and mathematics
 - Rules
 - To determine where widgets are placed and in what order
 - Which widgets are restricted in size
 - Which widgets can expand to take space
 - Mathematics
 - To give precise pixel values for position and size.
 - To produce a visually appealing order to widgets
 - Takes the form of algebraic equations and weighting systems

Layouts

- With these two things in place
 - A GUI program is capable of determining where every widget should be placed
 - And how much space it should be given
 - It allows the GUI program to scale up and scale down on a window resize operation.
 - Thus a GUI program can now take advantage of extra pixels should a user allocate them to the GUI window
 - All that's left to solve is how the widgets should be arranged

Layouts

- Arrangement of widgets
 - There is an infinite number of ways in which widgets can be arranged
 - Only a set few of these ways are used commonly throughout GUI programming
 - Why?

Layouts

- The human brain and visual system loves order and patterns
 - There is an infinite number of ways in which widgets can be arranged
 - Only a set few of these ways are used commonly throughout GUI programming
 - Why?

Layouts

- Patterns are easy to discern and remember
 - Horizontal line, vertical line, grid, border arrangement
 - Distinctive easy patterns
 - Because all potential users understand these patterns
 - Many applications will take advantage of them
 - Why?

Layouts

- Because it reduces the learning curve of GUIs significantly.
 - You are reusing existing knowledge of your users
 - Requires less mental effort to learn your application and interface.

Conclusion

- At this stage you now have a complete high level understanding of
 - How a window manager achieves window overlap
 - How a window manager determines what application owns an event
 - How an application will determine what widget owns an event
 - And subsequently processes it
 - And also how layouts work internally