



Computer Graphics 3: The graphics pipeline and graphics APIs

Programmer view of the graphics pipeline



- numerous ways that a user can interact with a graphics system
- For instance a painting program in the user sees menus and icons that represent possible actions
- We can also interact with the graphics pipeline through programming using an API

Graphics API



- There are many Graphics APIs
- The 2 most common are Open GL and DirectX and its 3d version Direct3D
- OpenGL is cross platform, open standard
- Originally Created by Silicon graphics, the responsibility for maintaining the standard has been turned over to the Khronos group

Graphics API

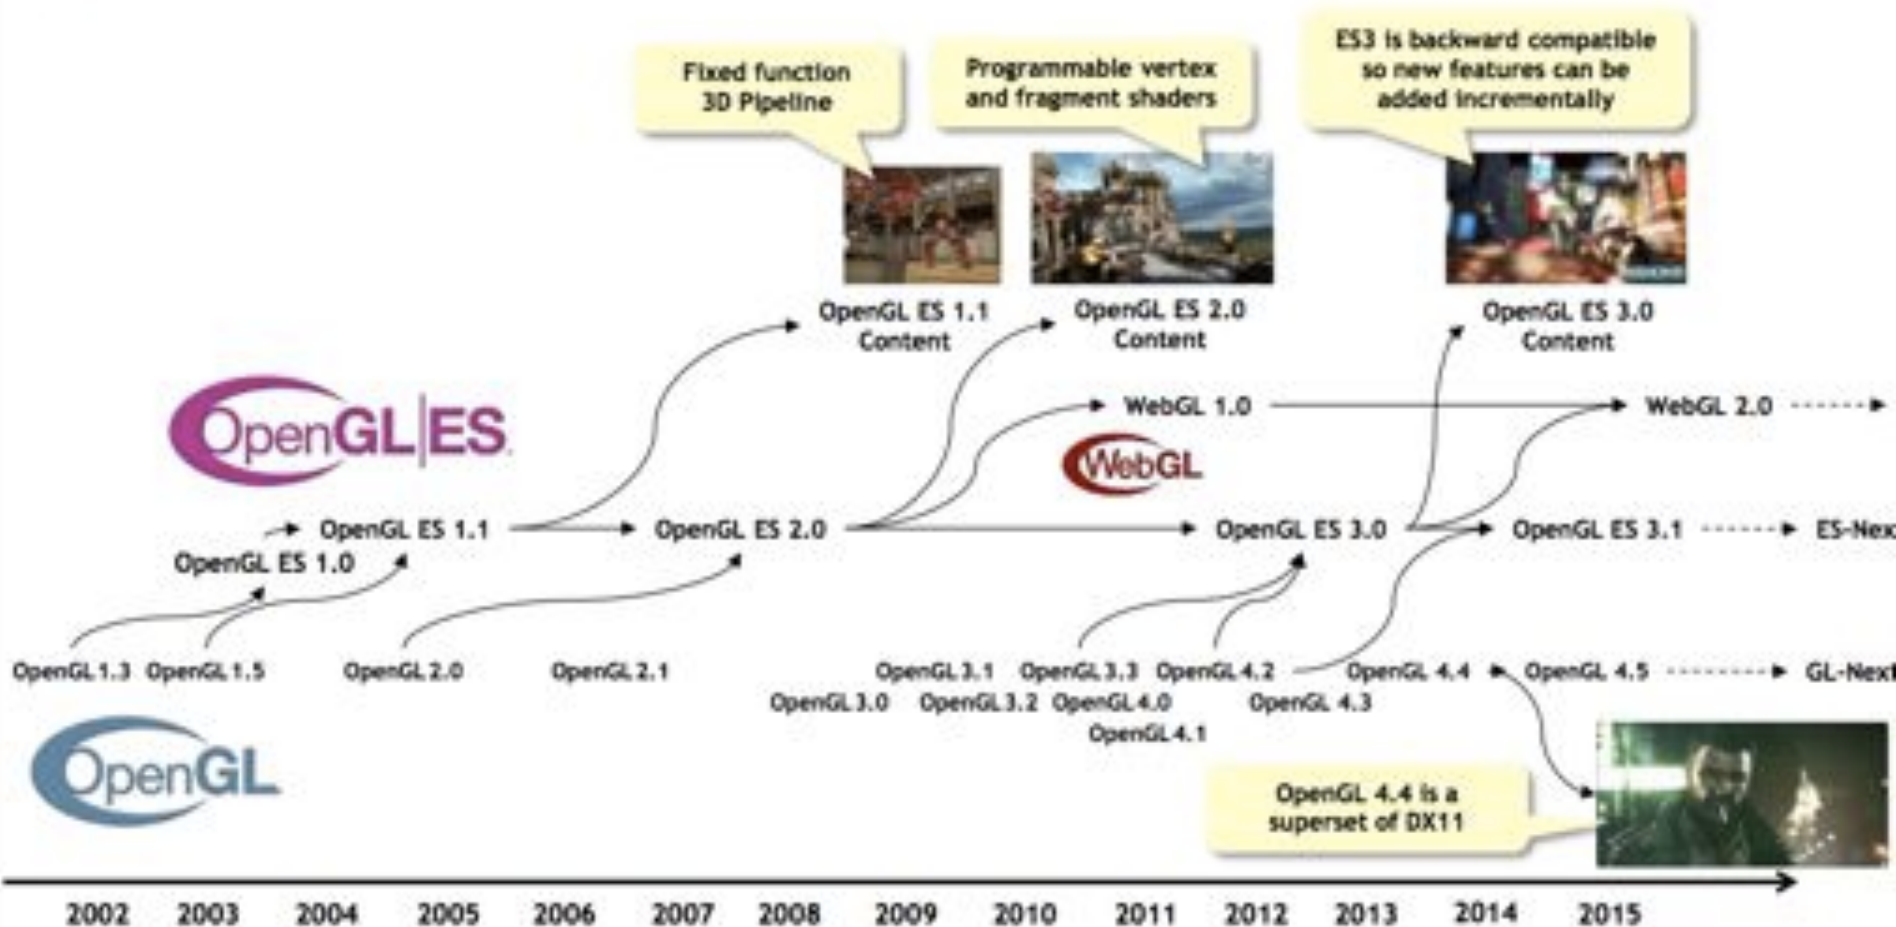


- There are many Graphics APIs
- The 2 most common are Open GL and DirectX and its 3d version Direct3D
- OpenGL is cross platform, open standard
- Originally Created by Silicon graphics, the responsibility for maintaining the standard has been turned over to the Khronos group

Graphics API



OpenGL 3D API Family Tree

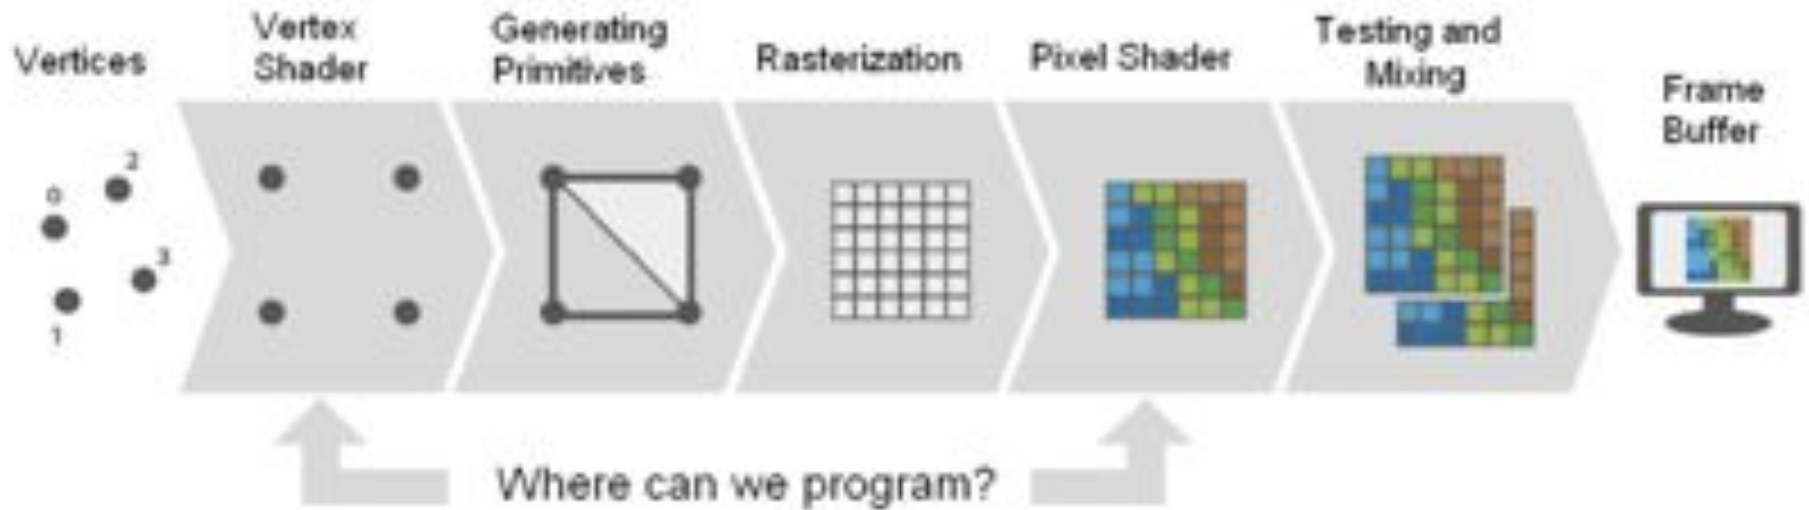


Outline of rendering pipeline



- rendering pipeline
- vertex processing
- primitive processing
- fragment processing
- summary

Outline of graphics pipeline



Rendering


- the process of generating an image given
 - a virtual camera
 - objects
 - light sources
- various techniques, e. g.
 - rasterization (topic of this course)
 - raytracing (topic of the course “Advanced Computer Graphics”)
- one of the major research topics in computer graphics
 - rendering
 - animation
 - geometry processing





Rasterization


- rendering algorithm for generating 2D images from 3D scenes
- transforming geometric primitives such as lines and polygons into raster image representations, i. e. pixels



[Akenine-Moeller et al.: Real-time Rendering]

Rasterization

- 3D objects are approximately represented by vertices (points), lines, polygons
- these primitives are processed to obtain a 2D image



[Akenine-Moeller]



Rendering Pipeline

- processing stages comprise the rendering pipeline (graphics pipeline)
- supported by commodity graphics hardware
 - GPU - graphics processing unit
 - computes stages of the rasterization-based rendering pipeline
- OpenGL and DirectX are software interfaces to graphics hardware
 - this course focuses on concepts of the rendering pipeline
 - this course assumes OpenGL in implementation-specific details



Outline

- introduction
- rendering pipeline
- vertex processing
- primitive processing
- fragment processing
- summary






Rendering Pipeline - Task



- 3D input
 - a virtual camera
 - position, orientation, focal length
 - objects
 - points (vertex / vertices), lines, polygons
 - geometry and material properties (position, normal, color, texture coordinates)
 - light sources
 - direction, position, color, intensity
 - textures (images)
- 2D output
 - per-pixel color values in the framebuffer

Rendering Pipeline / Some Functionality


- resolving visibility
- evaluating a lighting model
- computing shadows (not core functionality)
- applying textures



visibility



lighting model



shadow



texture

[Wright et al.: OpenGL SuperBible]

Rendering Pipeline

Main Stages

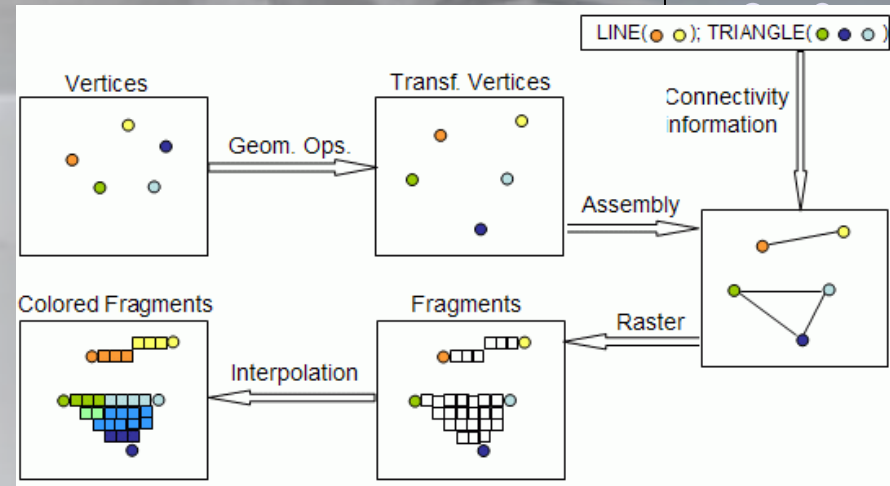


- vertex processing / geometry stage / vertex shader
 - processes all vertices independently in the same way
 - performs transformations per vertex, computes lighting per vertex
- geometry shader
 - generates, modifies, discards primitives
- primitive assembly and rasterization / rasterization stage
 - assembles primitives such as points, lines, triangles
 - converts primitives into a raster image
 - generates fragments / pixel candidates
 - fragment attributes are interpolated from vertices of a primitive
- fragment processing / fragment shader
 - processes all fragments independently in the same way
 - fragments are processed, discarded or stored in the framebuffer

Rendering Pipeline

Main Stages

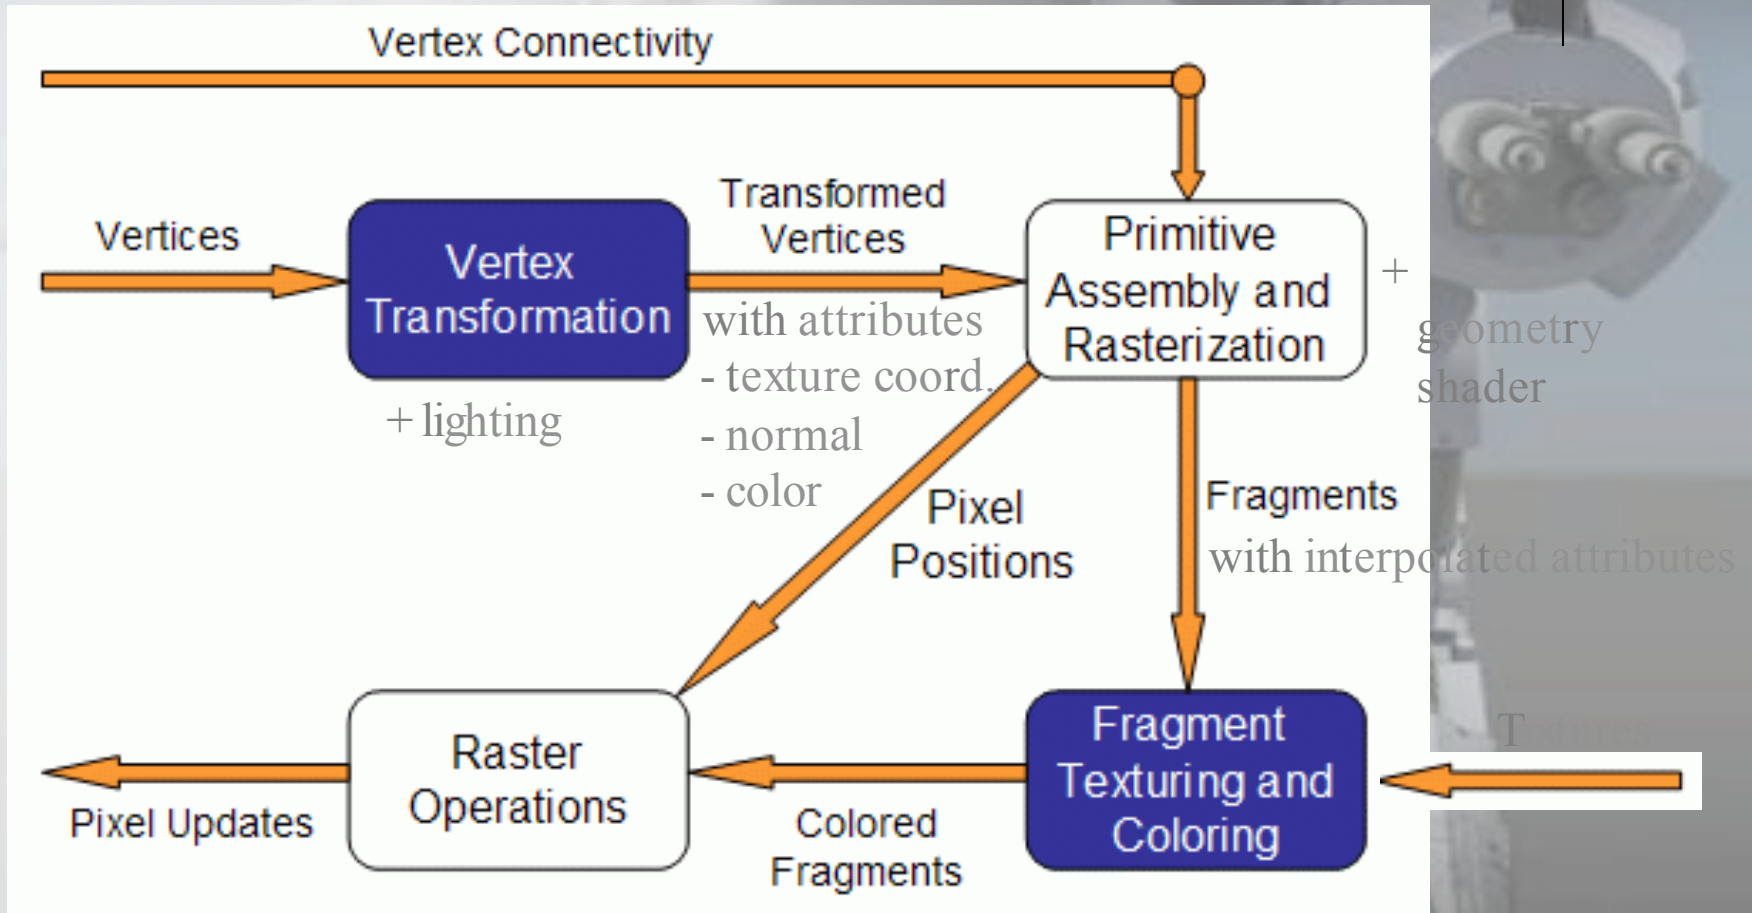
- vertex position transform
- lighting per vertex
- primitive assembly, combine vertices to lines, polygons
- rasterization, computes pixel positions affected by a primitive
- fragment generation with interpolated attributes, e. g. color
- fragment processing (not illustrated), fragment is discarded or used to update the pixel information in the framebuffer, more than one fragment can be processed per pixel position



[Lighthouse 3D]

Rendering Pipeline

Main Stages



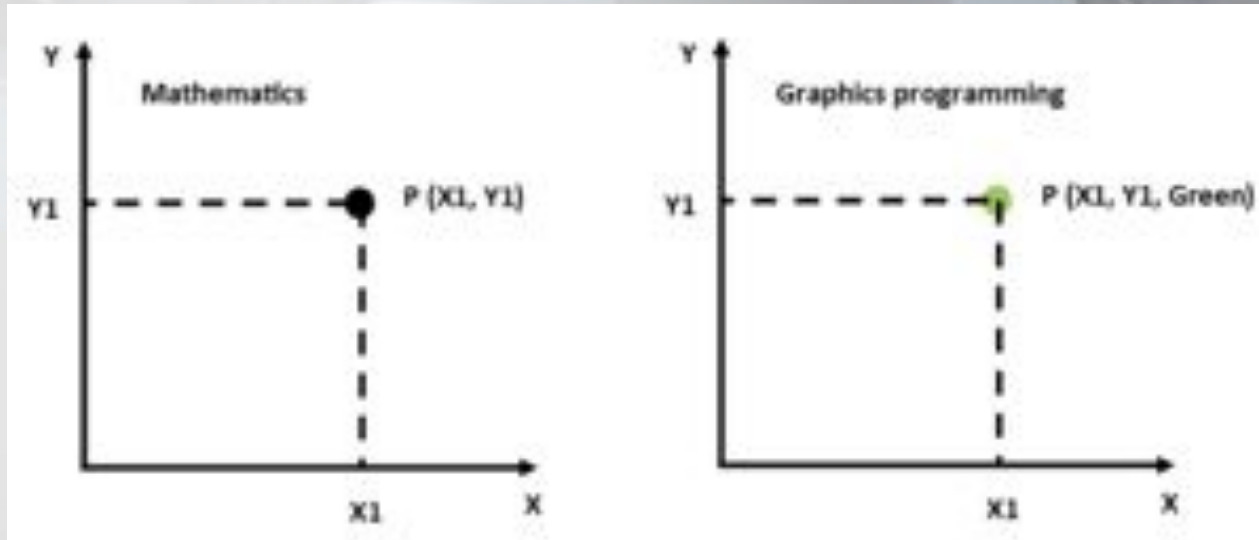
[Lighthouse 3D]

Outline

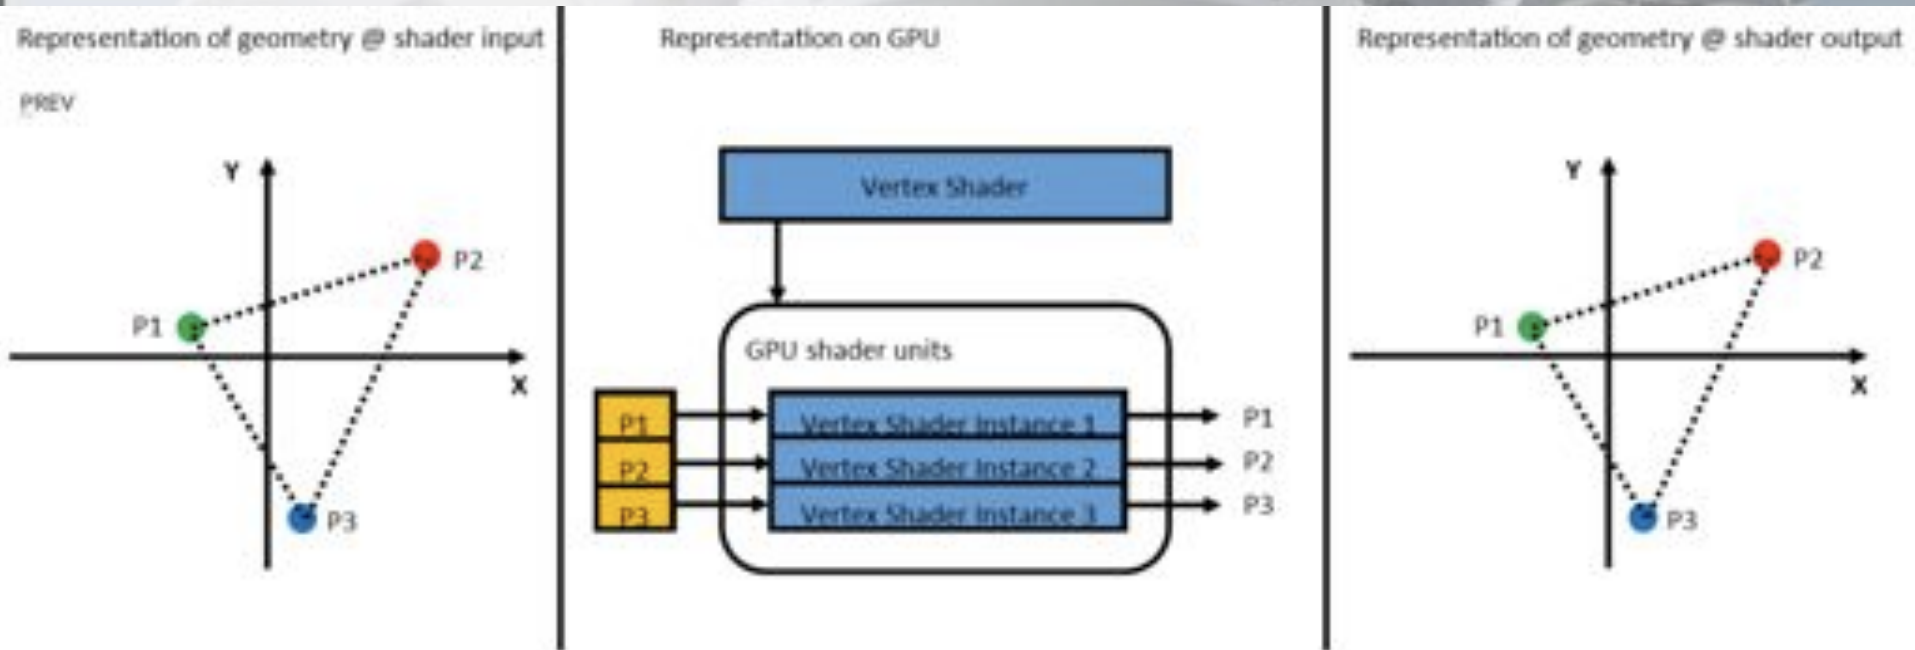
- introduction
- rendering pipeline
- **vertex processing**
- primitive processing
- fragment processing
- summary




Vertex



Shaders - Vertex Shaders



Primitive Assembly



Rasterization



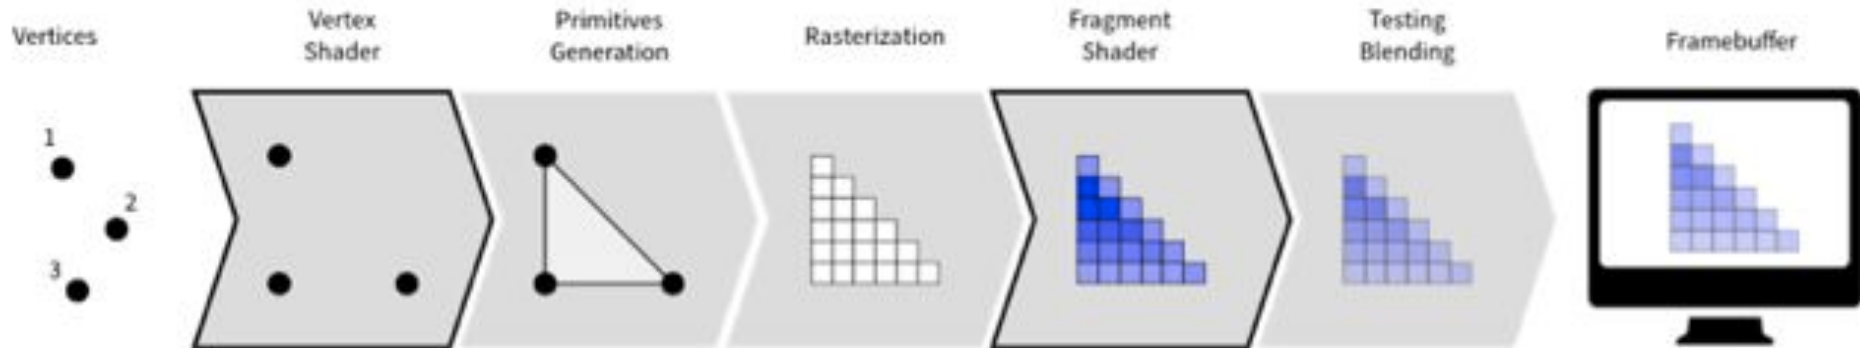


Figure 4: Rasterization of a primitive into 58 fragments

Fragment Shaders



Fragment Shaders



Vertex Processing (Geometry Stage)

- model transform
- view transform
- lighting
- projection transform
- clipping
- viewport transform



Model Transform

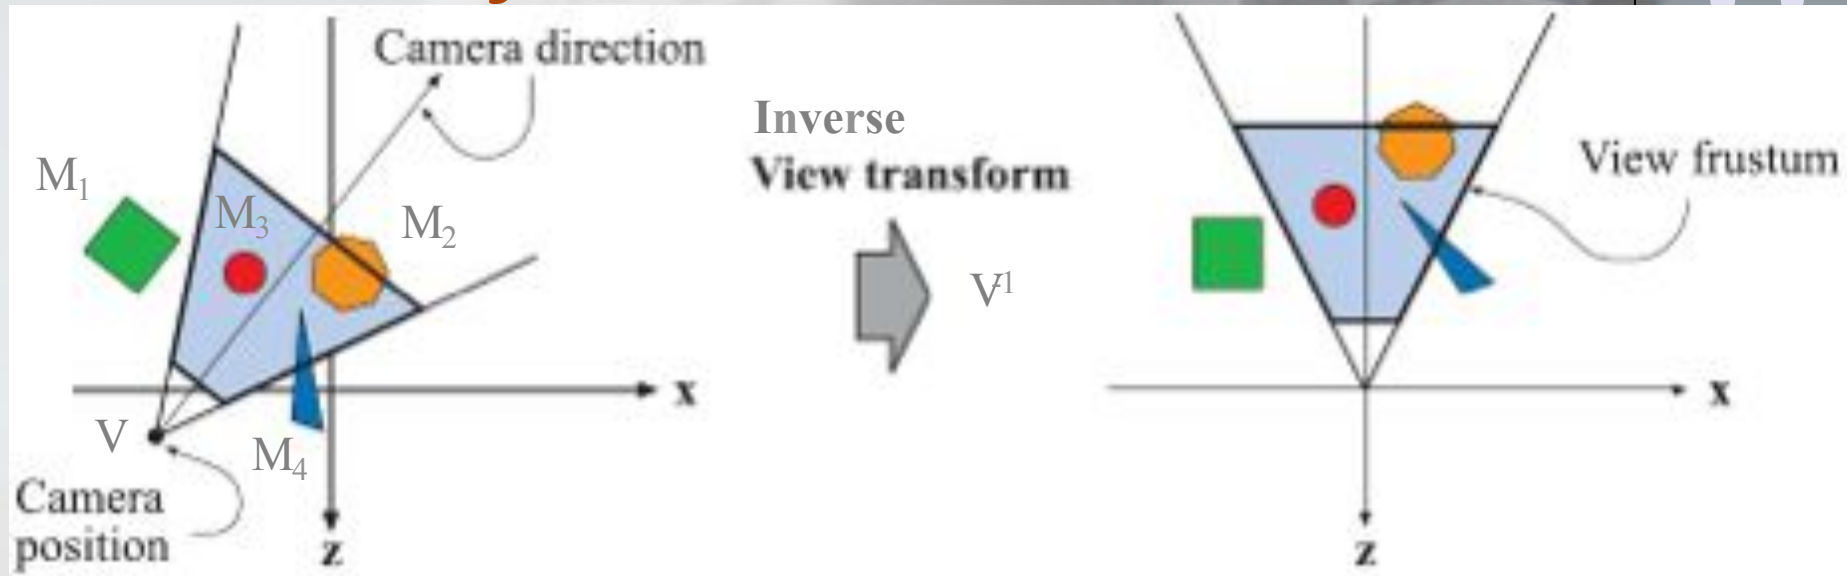
View Transform



- each object and the respective vertices are positioned, oriented, scaled in the scene with a model transform
- camera is positioned and oriented, represented by the view transform
- i. e., the inverse view transform is the transform that places the camera at the origin of the coordinate system, facing in the negative z-direction
- entire scene is transformed with the inverse view transform

Model Transform

View Transform



- M_1, M_2, M_3, M_4, V are matrices representing transformations
- M_1, M_2, M_3, M_4 are model transforms to place the objects in the scene
- V places and orientates the camera in space
 - V^{-1} transforms the camera to the origin looking along the negative z-axis
- model and view transforms are combined in the modelview transform
- the modelview transform $V^{-1}M_{1..4}$ is applied to the objects


Lighting

- interaction of light sources and surfaces is represented with a lighting / illumination model
- lighting computes color for each vertex
 - based on light source positions and properties
 - based on transformed position, transformed normal, and material properties of a vertex



Projection Transform

- P transforms the view volume to the canonical view volume
- the view volume depends on the camera properties
 - orthographic projection ③ cuboid
 - perspective projection ③ pyramidal frustum



- canonical view volume is a cube from $(-1, -1, -1)$ to $(1, 1, 1)$
- view volume is specified by near, far, left, right, bottom, top




Projection Transform

- view volume (cuboid or frustum) is transformed into a cube (canonical view volume)
- objects inside (and outside) the view volume are transformed accordingly
- orthographic
 - combination of translation and scaling
 - all objects are translated and scaled in the same way
- perspective
 - complex transformation
 - scaling factor depends on the distance of an object to the viewer
 - objects farther away from the camera appear smaller

Clipping

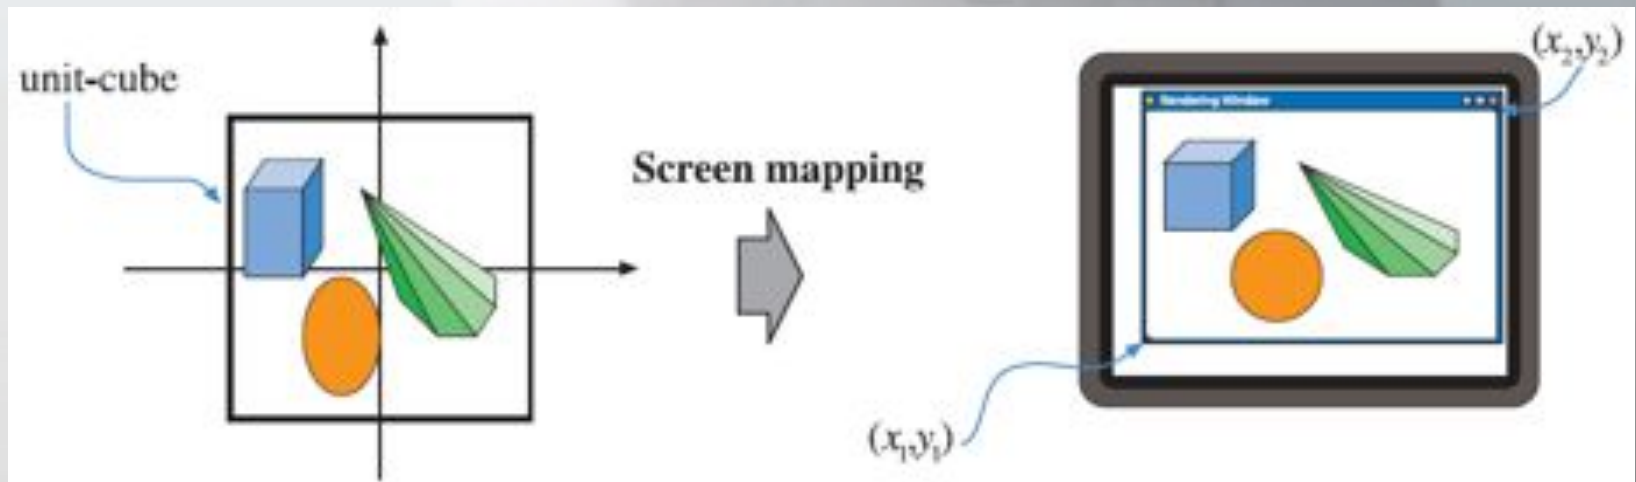
- primitives, that intersect the boundary of the view volume, are clipped
 - primitives, that are inside, are passed to the next processing stage
 - primitives, that are outside, are discarded
- clipping deletes and generates vertices and primitives



[Akenine-Moeller et al.: Real-time Rendering]

Viewport Transform / Screen Mapping

- projected primitive coordinates (x_p, y_p, z_p) are transformed to screen coordinates (x_s, y_s)
- screen coordinates together with depth value are window coordinates (x_s, y_s, z_w)



Viewport Transform / Screen Mapping



- (x_p, y_p) are translated and scaled from the range of $(-1, 1)$ to actual pixel positions (x_s, y_s) on the display
- z_p is generally translated and scaled from the range of $(-1, 1)$ to $(0, 1)$ for z_w
- screen coordinates (x_s, y_s) represent the pixel position of a fragment that is generated in a subsequent step
- z_w , the depth value, is an attribute of this fragment used for further processing



Vertex Processing - Summary

object space

»»» modelview transform

eye space / camera space

»»» lighting, projection

clip space / normalized
device coordinates

»»» clipping, viewport transform

window space



Vertex Processing - Summary

- input
 - vertices in object / model space
 - 3D positions
 - attributes such as normal, material properties, texture coords
- output
 - vertices in window space
 - 2D pixel positions
 - attributes such as normal, material properties, texture coords
 - additional or updated attributes such as
 - normalized depth (distance to the viewer)
 - color (result of the evaluation of the lighting model)

Outline

- introduction
- rendering pipeline
- vertex processing
- **primitive processing**
- fragment processing
- summary




Primitive Assembly / Rasterization



- primitive assembly
 - vertex information and connectivity information are
 - combined for further processing of points, lines and triangles
- geometry shader
 - change, delete, generate primitives
- rasterization
 - converts primitives into fragments
 - computes positions of screen pixels
 - that are affected by a primitive
 - generates a fragment for each affected pixel position
 - interpolates attributes from vertices to fragments

Rasterization

- input
 - vertices with attributes and connectivity information
 - attributes: color, depth, texture coordinates
- output
 - fragments with attributes
 - pixel position
 - interpolated color, depth, texture coordinates



Raster operations

- introduction
- rendering pipeline
- vertex processing
- primitive processing
- **fragment processing**
- summary



Raster operations: Fragment Processing



- fragment attributes are processed and tests are performed
 - fragment attributes are processed
 - fragments are discarded or
 - fragments pass a test and finally update the framebuffer
- processing and testing make use of
 - fragment attributes
 - textures
 - additional data that is available for each pixel position
 - depth buffer, color buffer, stencil buffer, accumulation buffer

Raster operations:Attribute Processing



- texture lookup
 - use texture coords to look up a texel (pixel of a texture image)
- texturing
 - combination of color and texel
- fog
 - adaptation of color based on fog color and depth value
- antialiasing
 - adaptation of alpha value (and color)
 - color has three components: red, green, blue
 - color is represented as a 4D vector (red, green, blue, alpha)



Raster operations: Tests

- scissor test
 - check if fragment is inside a specified rectangle
 - used for, e. g., masked rendering
- alpha test
 - check if the alpha value fulfills a certain requirement
 - comparison with a specified value
 - used for, e. g., transparency and billboarding
- stencil test
 - check if the stencil value in the framebuffer at the position of the fragment fulfills a certain requirement
 - comparison with a specified value
 - used for various rendering effects, e. g. masking, shadows



Raster operations: Depth Test

- depth test
 - compare depth value of the fragment and depth value of the framebuffer at the position of the fragment
 - used for resolving the visibility
 - if the depth value of the fragment is larger than the framebuffer depth value, the fragment is discarded
 - if the depth value of the fragment is smaller than the framebuffer depth value, the fragment passes and (potentially) overwrites the current color and depth values in the framebuffer

Raster operations: Blending / Merging



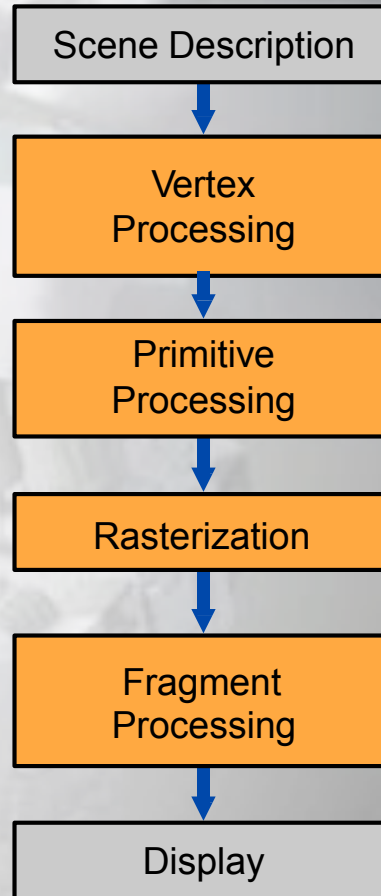
- blending
 - combines the fragment color with the framebuffer color at the position of the fragment
 - usually determined by the alpha values
 - resulting color (including alpha value) is used to update the framebuffer

Raster operations: Fragment Processing - Summary



- texture lookup
- texturing
- fog
- antialiasing
- scissor test
- alpha test
- stencil test
- depth test
- blending
- dithering
- logical operations

Rendering Pipeline - Summary





Rendering Pipeline - Summary

- primitives consist of vertices
- vertices have attributes (color, depth, texture coords)
- vertices are transformed and lit
- primitives are rasterized into fragments / pixel candidates with interpolated attributes
- fragments are processed using
 - their attributes such as color, depth, texture coordinates
 - texture data / image data
 - framebuffer data / data per pixel position (color, depth, stencil, accumulation)
- if a fragment passes all tests, it replaces the pixel data in the framebuffer